Sai Srinath Josyula

LinkedIn | Portfolio | srinath.5893@gmail.com

CAREER SUMMARY

Staff Software Engineer with over 9 years delivering backend platforms, distributed systems, and production operations at scale. Builds AI-native engineering programs: agentic automation, LLM orchestration, RAG, and Model Context Protocol (MCP) integrations that accelerate SDLC, on-call, and operations—with guardrails and human-in-the-loop design. Deep experience in Java and Spring microservices, APIs, streaming and batch data systems, and reliability at high throughput (including environments exceeding 1 million active users or transactions per second).

EDUCATION

Masters in Information Technology

Arizona State University, AZ · Jan 2017 – Dec 2018 · CGPA: 4/4

Related Coursework:

- Analyzing Big Data, Advance Data Analytics, Cloud Computing, Advanced SQL Techniques, System Administration of Unix, Software Engineering, Networking, Algorithms and Data Structures.

SKILLS AND ABILITIES

Languages: Java, Python, Scala, JavaScript

Web Technologies: HTML, XML, JavaScript, jQuery, ReactJS, Bootstrap, CSS

Big Data: Spark, MapReduce, Kafka, HBase, Hive, Oozie

API Related: Spring, ReactJS, GraphQL

AI, GenAI & LLMs: Generative AI, LLM integration & orchestration, RAG, embeddings, prompt and system design, tool-calling, evals, Copilot / GenAI product patterns, model evaluation & safety awareness, AI-assisted and agentic SDLC

Agentic & AI platforms: Agentic AI, AI agents and skills, multi-agent orchestration and workflow automation, Model Context Protocol (MCP)—first-party, Google Cloud MCP, and custom MCPs; Sentry, Camunda, Splunk, Jira, PostgreSQL integrations; HITL guardrails; Slack, GitHub, and M365 tool connectors; production AI delivery

AI-native dev & agent UIs: Claude (terminal/CLI and Claude Code), Claude Desktop, Cursor, Windsurf, Qodo — MCP-enabled workflows, multi-tool agent setups, and local/remote agent orchestration

Cloud, DevOps & more: Azure Blob Storage, Jenkins, XLR, Docker, AWS, Google Cloud, Maven, Kubernetes, Event Engine

WORK EXPERIENCE

Intuit Sept 2025 – Present

Staff Software Engineer

- **Agentic AI (on-call bot)** — Designed and built a production agentic investigation bot: an AI agent orchestrating Splunk, read-only DBs, GitHub, Slack, and Microsoft 365 via MCP for triage/RCA; owned prompts, tool flows, human-in-the-loop guardrails, and Claude Desktop/Cursor integration with secrets and multi-env Splunk.
- **AI-first SDLC** — Spearheaded an AI-first workspace where AI drives the SDLC end to end for Intuit merchant onboarding—agentic, model-native delivery on Claude Code and MCP, with agents in build/run/operate. Team shipped 40+ production AI agents/skills: Sentry triage, Camunda remediation, RAG, fraud, PR/code review, CR automation, Splunk investigations, security scanning, Jira/release/test—all agent-orchestrated on Splunk, Jira, Slack, GitHub, PostgreSQL. Multi-agent orchestration for autonomous on-call, zero-touch incident loops, Jira-to-code; AI as default interface for engineering and operations.
- **Kustomize / Kubernetes** — Refactored deployment configuration with Kustomize and a clear inheritance model so shared behavior is defined once and env-specific deltas stay small, reducing duplication, drift, and review overhead. Standardized layering across six environments improved readability and safe change velocity, preserved zero-downtime rollouts for five microservices.
- **Legacy reliability** — Improved a 15+ year legacy stack toward 99.99999% availability with 0 504s/500s, 0 release downtime, and 0 connection failures—removed chronic pager noise via scaling, traffic shaping, and controlled rollouts.
- **DeepTap-Spy** — Built DeepTap-Spy, a zero-code driver telemetry library unifying JDBC, Cassandra, Redis, Kafka, and AWS SDK v2; Micrometer-first plus interceptors elsewhere; mapped metrics into nine standard buckets; Prometheus, Wavefront, Splunk JSON, OTLP exporters.
- **Distributed locking** — Implemented a cache-agnostic locking framework (Redis) for Account Service (10M+ merchants, 100+ pods): composite keys, PII masking, four retry policies, lease watchdog, optional fencing tokens, and MDC-rich audit logging for production debuggability.

U.S. Bank May 2020 – Sept 2025

Senior Software Engineer

- Designed and implemented a highly scalable Ephemeral Search Service on Azure Blob Storage, enabling advanced search capabilities such as wildcards, range-based searches, exact matches, pagination, sorting, and limiting features. Established Spark batch pipelines to transfer data from Cassandra to blob storage, resulting in a cost reduction of approximately $3 million for the company by eliminating costly indexes, high computational resources needed for servers to run these indexes, and the maintenance needed for indexes and production issues. This service is currently used by US Bank to fulfill transaction search and wildcard text search requirements for mobile, online, and internal applications.
- Architected and fully automated the creation of a GraphQL microservice, eliminating manual intervention by automating GraphQL schema file generation, GIT project creation, code generation, Jenkins job creation, Postman collection generation, and Karate test suite creation. Enabled consumers to deploy a GraphQL microservice in 5 minutes by inputting Cassandra keyspace and table name. Currently, over 500 GraphQL microservices are running in production.
- Implemented robust security measures, including JWT generation and validation, to secure API endpoints, enhancing the overall security posture of the services.
- Developed custom circuit breaker logic to significantly enhance the resilience and reliability of microservices, currently utilized across multiple APIs.
- Built an internal ReactJS application, leveraging Apollo GraphQL and React Router for seamless integration with backend services, improving internal workflows and efficiency.
- Created utilities for seamless data integration with Kafka, enabling efficient message processing across various payload structures, enhancing data flow and processing capabilities.
- Applied industry-standard design patterns and best practices to ensure high code quality and maintainability, contributing to the development of robust and scalable frameworks.
- Utilized Maven for project management and build automation, streamlining the development process and improving productivity.

Impetus Jan 2019 – May 2020

Software Engineer

- Contributed to the Batch Processing team, developing complex MapReduce and Spark applications using Java, Scala, and Python to process large volumes of data efficiently.
- Extracted data from source Hive tables, applied Drools rules, and loaded the processed data into HBase tables. Developed shell scripts for archiving batch data and Oozie scripts for job scheduling.
- Developed a standalone Java application to read data from Hive and HBase tables and push it to Kafka in the required format. The application is triggered using the Event Engine job scheduler. Also developed Kafka consumers for data consumption and testing.
- Proficient in using build tools like Maven and Gradle and experienced in using XLR to manage the entire Software Development Life Cycle.

EdgeVerve Dec 2014 – Dec 2016

Software Engineer

- Built RESTful web services using Java and Spring, with Cassandra as the backend database, ensuring efficient data handling and retrieval.
- Developed enhancements for Finacle Treasury, including migrating application server configurations to the database, improving system manageability and performance.

MASTERS PROJECTS

Twitter-Review Analysis Application

Developed an application to provide review analysis of places, food, events, etc., through charts. Utilized Django for the user interface, deployed the application on Google App Engine, stored Tweets in BigQuery, and visualized data using Google Data Studio.

Technologies: Python, Django, Google App Engine, BigQuery, Google Data Studio.

Meetup Portal for ASU Polytechnic Campus

Created a portal enabling students and faculty to organize meetings and share knowledge.

Technologies: Java, JDBC, Servlets, HTML, CSS, JavaScript.

Photo-voltaic Module Certification Portal

Built a fully functional portal for Photo-voltaic Module certification using Python and the Django framework. Employed MSSQL as the database and implemented stored procedures, cursors, triggers, views, and dynamic SQL.

Technologies: Python, Django, MSSQL.